

Systems Security

COMSM1500

Access Control

bristol.ac.uk



Access Control

- Physical access control
 - Physical enforcement to access to areas
- Digital access control
 - Restrict access to resources and interactions



Subjects and Objects


- Object
 - Passive entity that contains information
 - e.g. file, record, memory location etc.
- Access
 - Ability to perform an action on/interact with an object
 - Flow of information between a subject and an object
- Subject
 - Active entity requesting access to an object or data within an object
 - Different subjects have different access levels
 - e.g. users, processes etc.

Access Control

- Concerned with **authorisation**: what a subject is allowed to do
- **Mediates** a subject access to object
- **Enforces a security policy**, limiting which actions are allowed

Complete mediation

- That is our aim
- **Trusted computing base**: all hardware and software that is responsible for enforcing the policy
 - For example, OS kernel, all trusted processes, and the PC hardware. A fault in one can compromise security (we have seen many examples).
- Complete mediation is important (if it can be subverted, then use is limited)
- **Reference monitor** concept (all access through it)



Protection state

- **Security context:** security identity information used to inform authorisation decision.
 - For example, UID and GID associated with a process
- The **protection state** is made up of the security sensitive actions every entity is able to do
 - For example, when a user change identity the protection state has changed
- **Transition** between protection state need to be tightly controlled (e.g. setgid)



Access Control Matrix

- Simplest way to represent the protection state of a system
- A table representing every subjects and objects, and the permitted type of actions between them

	File 1	File 2	File 3
User A	Read, Write, Own	Read, Write	Read, Write, Own
User B	Append	Read, Write, Own	Read, Write

What's the fundamental issue here?



Access Control Matrix

- Could in theory express any possible access control policy
- Only useful in theory...
- In any practical systems there is too many subjects and objects
- ... but any protection state can be expressed this way

	File 1	File 2	File 3
User A	Read, Write, Own	Read, Write	Read, Write, Own
User B	Append	Read, Write, Own	Read, Write

Access control policy


- Express formally or informally what is allowed
- Generally express confidentiality/integrity requirements

	File 1	File 2	File 3
User A	Read, Write, Own	Read, Write	Read, Write, Own
User B	Append	Read, Write, Own	Read, Write


Mechanisms and Models

- A **mechanism** how the policy is enforced (think of software/hardware implementation)
 - e.g. SELinux, AppArmor on Linux distributions
- A **model** how to represent policies
 - e.g. access control matrix (this is not a good idea)

Access Control List



Capabilities




Access Control List/Capabilities

- Access Control List
 - Object centric
- Capabilities
 - Subject centric

Discretionary / Mandatory

- Discretionary: owner of a resource defines/delegates right
 - Requires the notion of owner
 - Most systems
 - Resource owner define associated policy
- Mandatory: policies centrally set by an administrator
 - Notion of owner optional
 - Well suited when an organization own the data (e.g. military government)

Access Control List




Access Control List

- On Windows work more or less like this (virtually unlimited list)
- On UNIX: 3 entries Owner, Group, World
- By default owner=creator
- Inheritance rules (objects inherit rules of containers)
 - e.g. files inherit folders permission
- Owner can **revoke** access
- DAC

UNIX

```
~/syssec/2017$ ls -l
-rw-r--r-- csxdb cosc ... accesscontrol.pdf
-rw-r--r-- csxdb cosc ... systems.pdf
-rw----- csxdb cosc ... exam.pdf
-rwxr-x--- csxdb cosc ... slides.sh
```



owner	group
-------	-------


owner	group	other
-------	-------	-------

read	write	exec
------	-------	------

Role-based Access Control

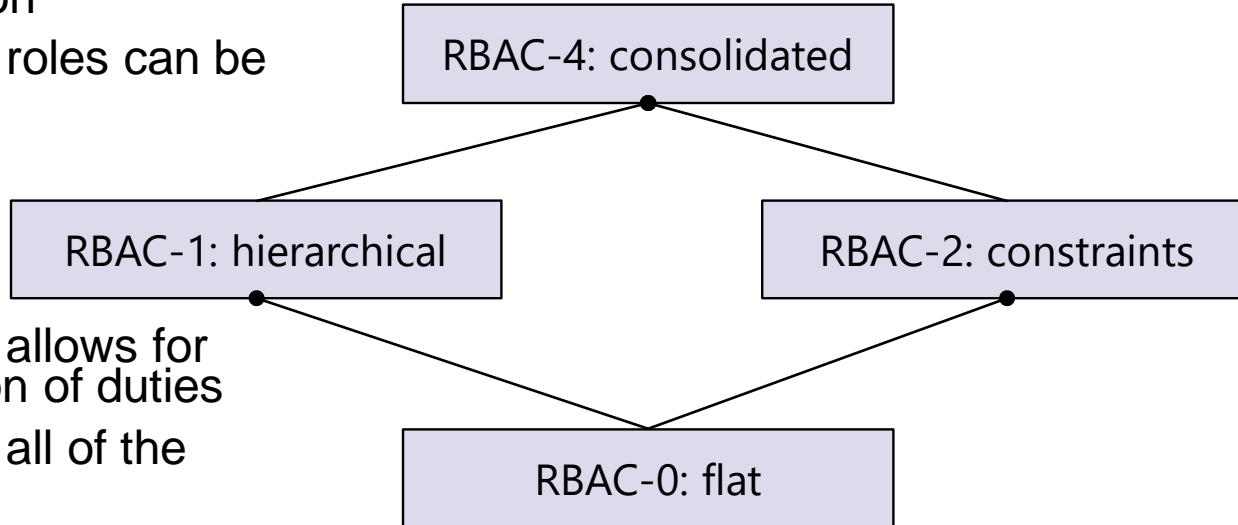


Groups / Roles



RBAC (NIST)

- RBAC 0: users have a role, role have permission
- RBAC 1: roles can be inherited



- RBAC 2: allows for separation of duties
- RBAC 4: all of the above

Information Flow Control

bristol.ac.uk



IFC

- read: receiving information from another entity
- write: sending information from another entity
- Concept emerges from US military

IFC- Bell Lapadula

Top-secret (2)
Secret (1)
Unclassified (0)

IFC- Bell Lapadula

Top-secret (2)
Secret (1)
Unclassified (0)

- read: $f_s(s) \geq f_o(o)$
–no read up

IFC- Bell Lapadula

Top-secret (2)
Secret (1)
Unclassified (0)

- read: $f_s(s) \geq f_o(o)$
–no read up
- write: $f_s(s) \leq f_o(o)$
–no write down

IFC- Bell Lapadula

Top-secret (2)
Secret (1)
Unclassified (0)

- read: $f_s(s) \geq f_o(o)$
–no read up
- write: $f_s(s) \leq f_o(o)$
–no write down
- read/write: $f_s(s) == f_o(o)$

Problem with this naïve approach?



IFC - Bell Lapadula

Top-secret (2)
Secret (1)
Unclassified (0)

- Need to add categories! (DAC)
- (c, s)
 - c classification
 - s categories
- $f_s(s)$ dominates $f_o(o)$
 - if $c_s \geq c_o$ MAC
 - e.g. top-secret \geq secret
 - if $s_s \supseteq s_o$ DAC
 - e.g. {Iraq, Syria} \supseteq {Iraq}


IFC- Bell Lapadula

Top-secret (2)
Secret (1)
Unclassified (0)

- read: $f_s(s)$ dominates $f_o(o)$
 - no read up
- write: $f_o(o)$ dominates $f_s(s)$
 - no write down
- read/write: $f_s(s) == f_o(o)$
- SELinux variant
 - Level per category

IFC - Bell Lapadula

- Categories have a partial order.
- Form a lattice



IFC - Biba Model (similar idea for integrity)

Fact (2)
Belief (1)
Rumor (0)

IFC - Biba Model (similar idea for integrity)

Fact (2)
Belief (1)
Rumor (0)

- read: $f_s(s) \leq f_o(o)$
–no read down

IFC - Biba Model (similar idea for integrity)

Fact (2)
Belief (1)
Rumor (0)

- read: $f_s(s) \leq f_o(o)$
–no read down
- write: $f_s(s) \geq f_o(o)$
–no write up

IFC - Biba Model (similar idea for integrity)

Fact (2)
Belief (1)
Rumor (0)

- read: $f_s(s) \leq f_o(o)$
–no read down
- write: $f_s(s) \geq f_o(o)$
–no write up
- read/write: $f_s(s) == f_o(o)$

IFC - Biba Model (similar idea for integrity)

Fact (2)
Belief (1)
Rumor (0)

- read: $f_s(s) \leq f_o(o)$
–no read down
- write: $f_s(s) \geq f_o(o)$
–no write up
- read/write: $f_s(s) == f_o(o)$
- Similarly we can add DAC categories

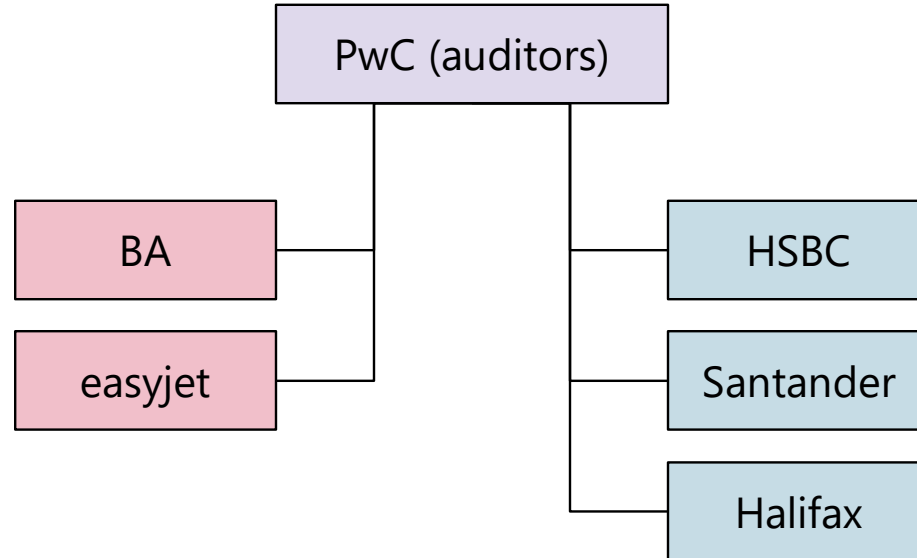
Brewer and Nash Model

a.k.a Chinese Wall

bristol.ac.uk



Chinese Wall Model




Chinese Wall Model

O = objects (clients' files)

C = clients

Col(o) = Conflict of Interest Group

x(o): owner, y(o): conflict class Col(o) – x(o)



Chinese Wall Model

- Define $n(s, o)$: has s ever read o ?

Chinese Wall Model

- Define $n(s, o)$: has s ever read o ?
- s can read o if:
 - For all o' such that $n(s, o') = \text{true}$, $y(o) = y(o')$ or $y(o) \not\equiv x(o')$
 - o is from the same company as previously read objects or o belong to a different conflict of interest class

Chinese Wall Model

- Define $n(s, o)$: has s ever read o ?
- s can read o if:
 - For all o' such that $n(s, o') = \text{true}$, $y(o) = y(o')$ or $y(o) \not\in x(o')$
 - o is from the same company as previously read objects or o belong to a different conflict of interest class
- s can write o if:
 - s can read o
 - $\nexists o'$ such that $n(s, o') = \text{true}$ with $y(o) \neq y(o')$ unless $x(o') = \emptyset$
 - You can't write client data if reading anyone else data, unless other data have been sanitized

Chinese Wall Model

- Define $n(s, o)$: has s ever read o ?
- s can read o if:
 - For all o' such that $n(s, o') = \text{true}$, $y(o) = y(o')$ or $y(o) \not\in x(o')$
 - o is from the same company as previously read objects or o belong to a different conflict of interest class
- s can write o if:
 - s can read o
 - $\nexists o'$ such that $n(s, o') = \text{true}$ with $y(o) \neq y(o')$ unless $x(o') = \emptyset$
 - You can't write client data if reading anyone else data, unless other data have been sanitized
 - Need sanitization/declassification policy
 - Simple example owner can declassify objects they own
 - Authority may declassify audit report etc.

Clark-Wilson Model

Integrity

bristol.ac.uk



Clark-Wilson Model

- Bank-like context
- Integrity defined as a set of constraints
 - Data in a valid state when it meet constraints
 - Example
 - YB: Yesterday Balance, TD: today deposit, TW: today withdrawal, TB: today balance
 - $TB = YB + TD - TW$
- Well formed transactions
 - Move system to one consistent state to another
- Require separation of duty
 - Who certify transactions
 - Who verify constraints


Clark-Wilson Model



constrained
data item



unconstrained
data item



transformation
procedure



integrity verification
procedure

Clark-Wilson Model

CDI

constrained
data item

UDI

unconstrained
data item

TP


transformation
procedure

IVP

integrity verification
procedure

- Permission as a triple
 - (user, TP, {CDIs})
- e.g. (bob, transfer, hisaccount)
- transfer(CDI:account, UDI:accountref)

Clark-Wilson Model



Certification rules

- C1: CDI state validated by IVP
- C2: TPs preserve valid state
- C3: separation of duty
- C4: TPs write to log
- C5: UDIs are validated by TPs

Enforcement rules

- E1: Only TPs change CDIs
- E2: TPs require authorisation
- E3: All users are authenticated
- E4: only authorised people can change permissions

Summary

- Many models for different objectives
- ... and there is many more
 - attribute-based access control
 - provenance-based access control
 - etc..
- Different model for different objectives (in theory you could obtain the same access control matrix from two different models)
- They have led to programming pattern
- In practice hard to implement
 - Many secure OS attempts (yet nothing widely available)
 - Real world is full of exception and edge cases
 - May get in the way of business

Thank you

Office MVB 3.26

bristol.ac.uk

