

# Systems Security

## COMSM1500

# Race condition vulnerabilities



# Plan

- Race condition
- Example of race condition vulnerabilities
  - access system call
  - How (not to) implement an OS reference monitor
  - dirty cow

# What is a race condition?



# What is a race condition?



# What is a race condition?



# What is a race condition?



Resource  
10



# What is a race condition?



Read




Resource  
10

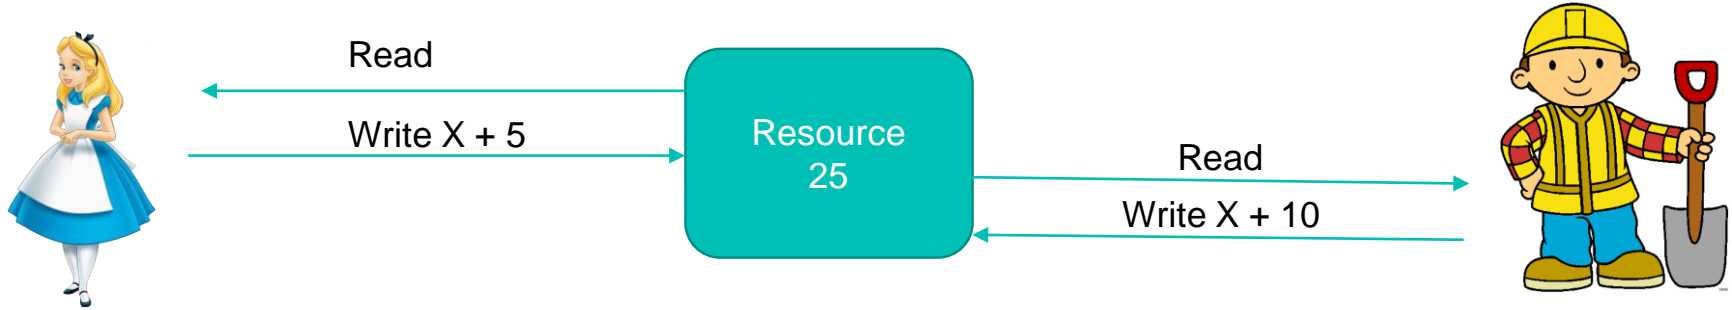




# What is a race condition?



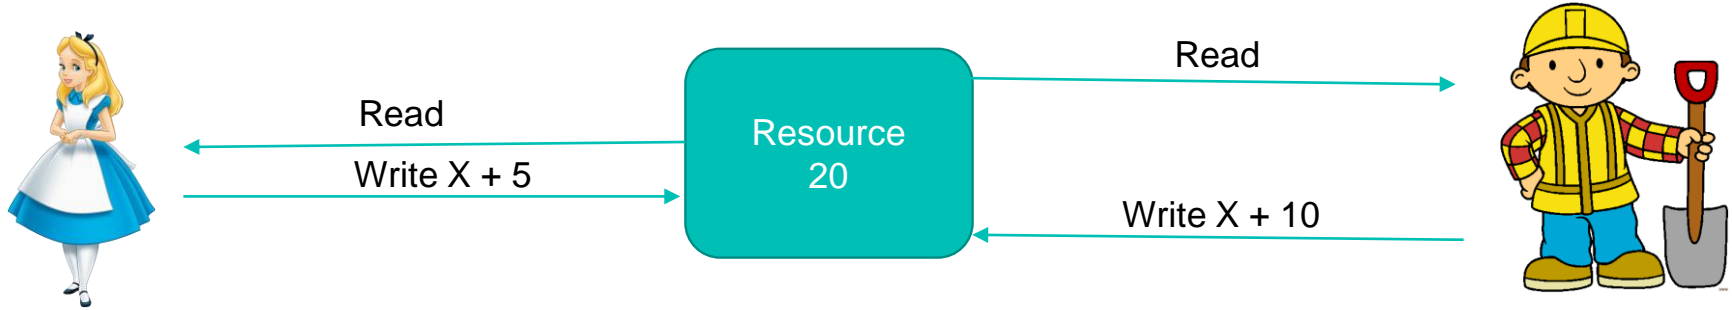
# What is a race condition?



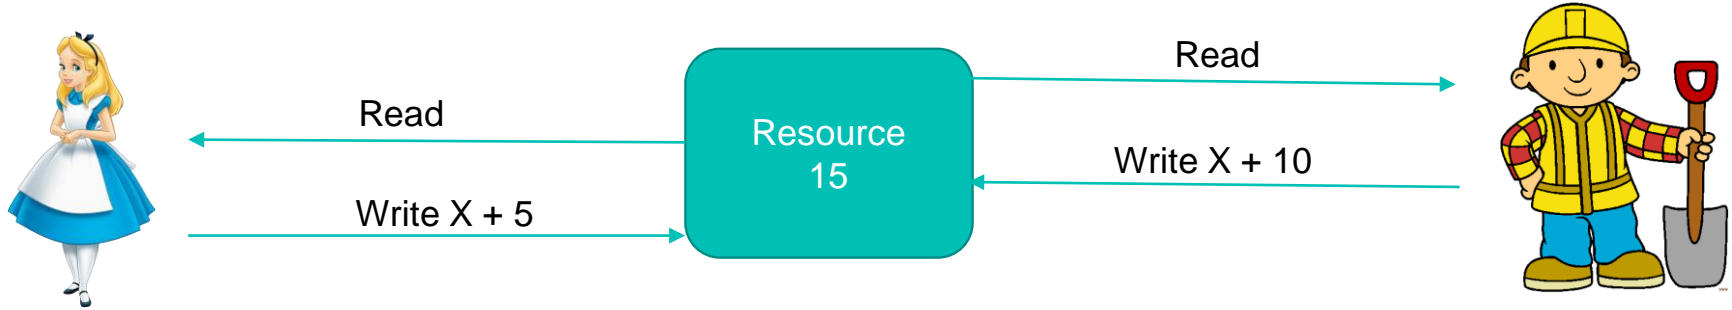
# What is a race condition?

- All good here!

# What is a race condition?



# What is a race condition?



# What is a race condition?

```
fcn withdraw(amount)
    balance = getBalance()
    if (balance > amount)
        balance = balance – amount
        saveBalance(balance)
    else
        print “not enough fund”
```

# What is a race condition?

- Race condition:
  - System where the output is dependent on the sequence or timing of events
- Many possible permutations
- ... you have seen this before
- It is a well understood problem, with solutions
  - Synchronization
  - Transactions
  - etc...

# Race condition vulnerabilities





# Example 1: access system call

```
if(access("tmp/X", W_OK)) {  
    f = open("tmp/X");  
    write_to_file(f);  
} else {  
    printf("You do not own the file");  
}
```

# Example 1: access system call

```
if(access("tmp/X", W_OK)) {  
    f = open("tmp/X");  
    write_to_file(f);  
} else {  
    printf("You do not own the  
    file");  
}
```

- setuid: root
- want to make sure the “real” user own the file

# Example 1: access system call

```
if(access("tmp/X", W_OK)) {  
    f = open("tmp/X");  
    write_to_file(f);  
} else {  
    printf("You do not own the  
    file");  
}
```

- setuid: root
- want to make sure the “real” user own the file
- access return either or not the operation is permitted to current user

# How can this be exploited?



# Example 1: access system call

```
if(access("tmp/X", W_OK)) {  
    f = open("tmp/X");  
    write_to_file(f);  
} else {  
    printf("You do not own the  
    file");  
}
```

- Path hard coded
  - Program will only write to /tmp/X

# Example 1: access system call

```
if(access("tmp/X", W_OK)) {1  
    f = open("tmp/X");2  
    write_to_file(f);  
} else {  
    printf("You do not own the  
    file");  
}
```

- Path hard coded
  - Program will only write to /tmp/X
- Symbolic Link
  - /tmp/X -> /etc/config
- 1: fail
- 2: success

# Example 1: access system call

```
if(access("tmp/X", W_OK)) {1  
    f = open("tmp/X");2  
    write_to_file(f);  
} else {  
    printf("You do not own the  
    file");  
}
```

- Exploited race condition
- Changed value between check and use
- time of check to time of use  
– TOCTOU

# access man

**Warning:** Using `access()` to check if a user is authorized to, for example, open a file before actually doing so using `open`(2) creates a security hole, because the user might exploit the short time interval between checking and opening the file to manipulate it. **For this reason, the use of this system call should be avoided.** (In the example just described, a safer alternative would be to temporarily switch the process's effective user ID to the real ID and then call `open`(2).)



# access man

Homework/exam question:  
access system call is vulnerable  
to race condition. Explain how it  
can be exploited.

**Warning:** Using `access()` to check if a user is authorized to, for example, open a file before actually doing so using `open(2)` creates a security hole, because the user might exploit the short time interval between checking and opening the file to manipulate it. **For this reason, the use of this system call should be avoided.** (In the example just described, a safer alternative would be to temporarily switch the process's effective user ID to the real ID and then call `open(2)`.)

## Example 2: implementing OS reference monitor

# What is a reference monitor?



# What is a reference monitor?

Seen in access control lecture  
on October 30<sup>th</sup>!

[bristol.ac.uk](http://bristol.ac.uk)

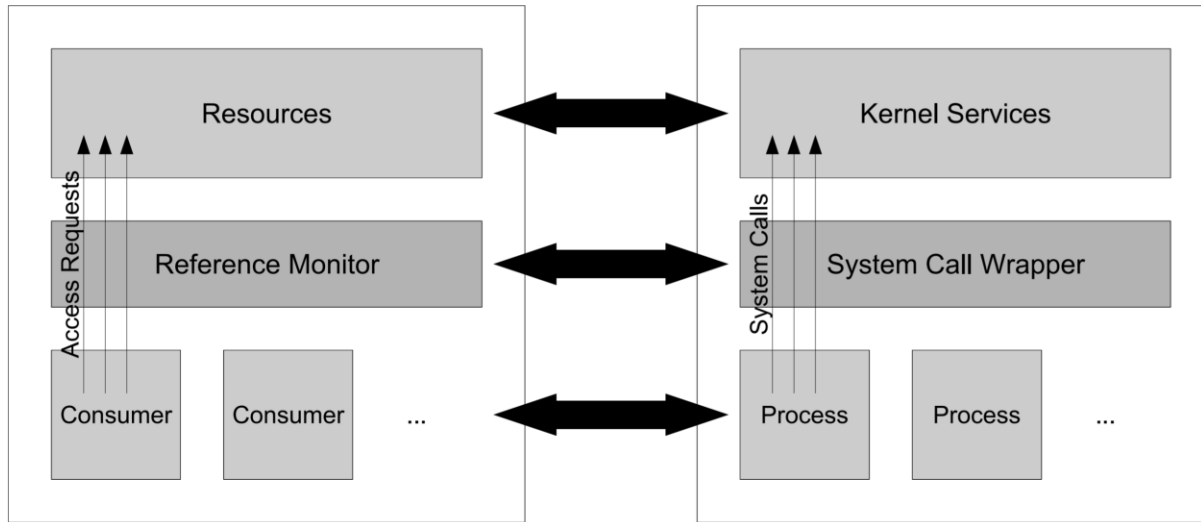


## Example 2: implementing OS reference monitor


- a **reference monitor** is a secure, always-used and fully-testable module that controls all software access to data objects or devices

# Example 2: implementing OS reference monitor


- a **reference monitor** is a secure, always-used and fully-testable module that controls all software access to data objects or devices



## Example 2: implementing OS reference monitor




## Example 2: implementing OS reference monitor



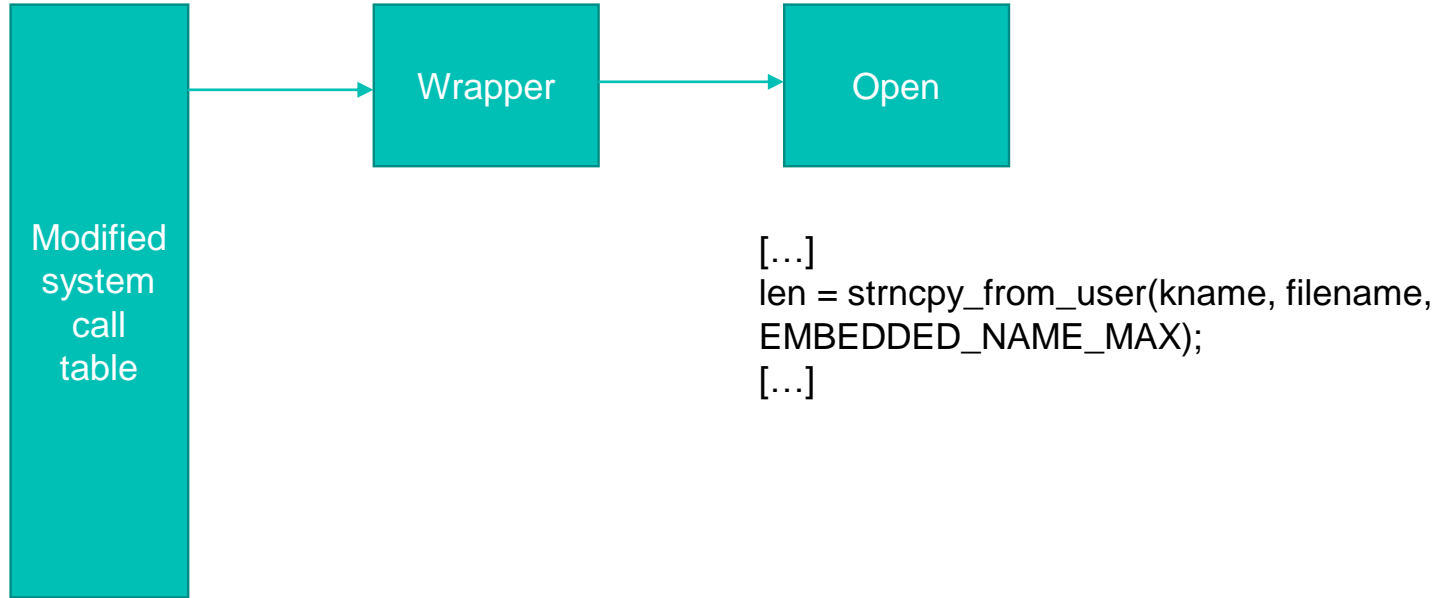


## Example 2: implementing OS reference monitor

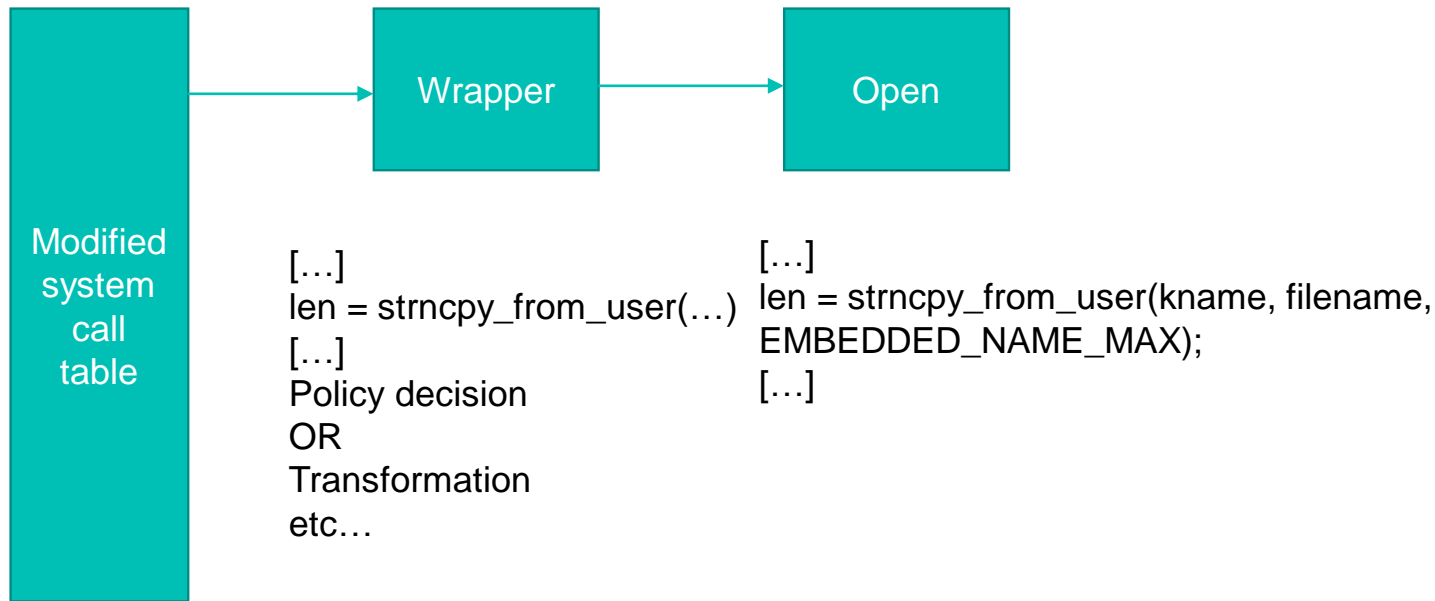


e.g. LOMAC

## Example 2: implementing OS reference monitor



## Example 2: implementing OS reference monitor



# Problem?



## Example 2: implementing OS reference monitor

- Wrapper and syscall work on two different copy of the buffer!
- User space controlled buffer
  - Can be controlled by an attacker
- The value checked to enforce policy!=value seen by syscall

## Example 2: implementing OS reference monitor

- Wrapper and syscall work on two different copy of the buffer!
- User space controlled buffer
  - Can be controlled by an attacker
- The value checked to enforce policy!=value seen by syscall
- Check LSM paper for how to properly implement a reference monitor!

## Example 2: implementing OS reference monitor

- Wrapper and syscall work on two different copy of the buffer!
- User space controlled buffer
  - Can be controlled by an attacker
- The value checked to enforce policy!=value seen by syscall
- Check LSM paper for how to properly implement a reference monitor!
- Check Robert Watson's paper for “real” exploit to syscall wrapper

## Example 2: implementing OS reference monitor

- Wrapper and syscall work on two different copy of the buffer!
- User space controlled buffer
  - Can be controlled by an attacker
- The value checked to enforce policy!=value seen by syscall

Homework/exam question:  
Explain why system call wrapper  
should not be used to implement  
access control.




## Example 3: dirty cow




## Example 3: dirty cow

- CVE-2016-5195
- Linux vulnerability that could be exploited to gain root access
- Concurrency issue relating to how memory is managed

## Example 3: dirty cow




Process memory




Physical memory

## Example 3: dirty cow




Process memory

Translation  
Table




Physical memory

## Example 3: dirty cow



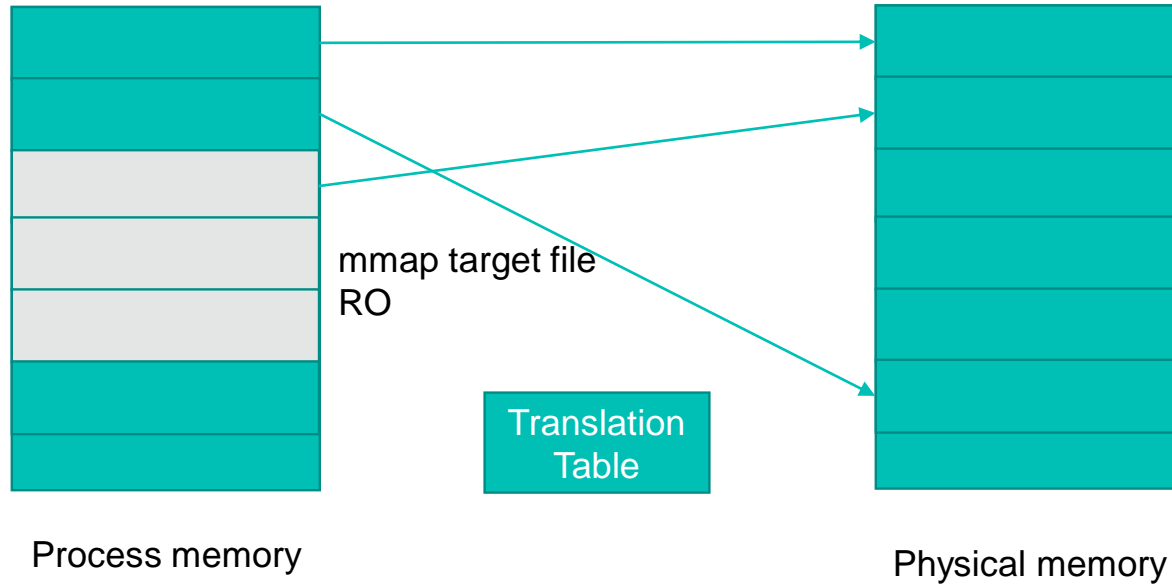
# Example 3: dirty cow



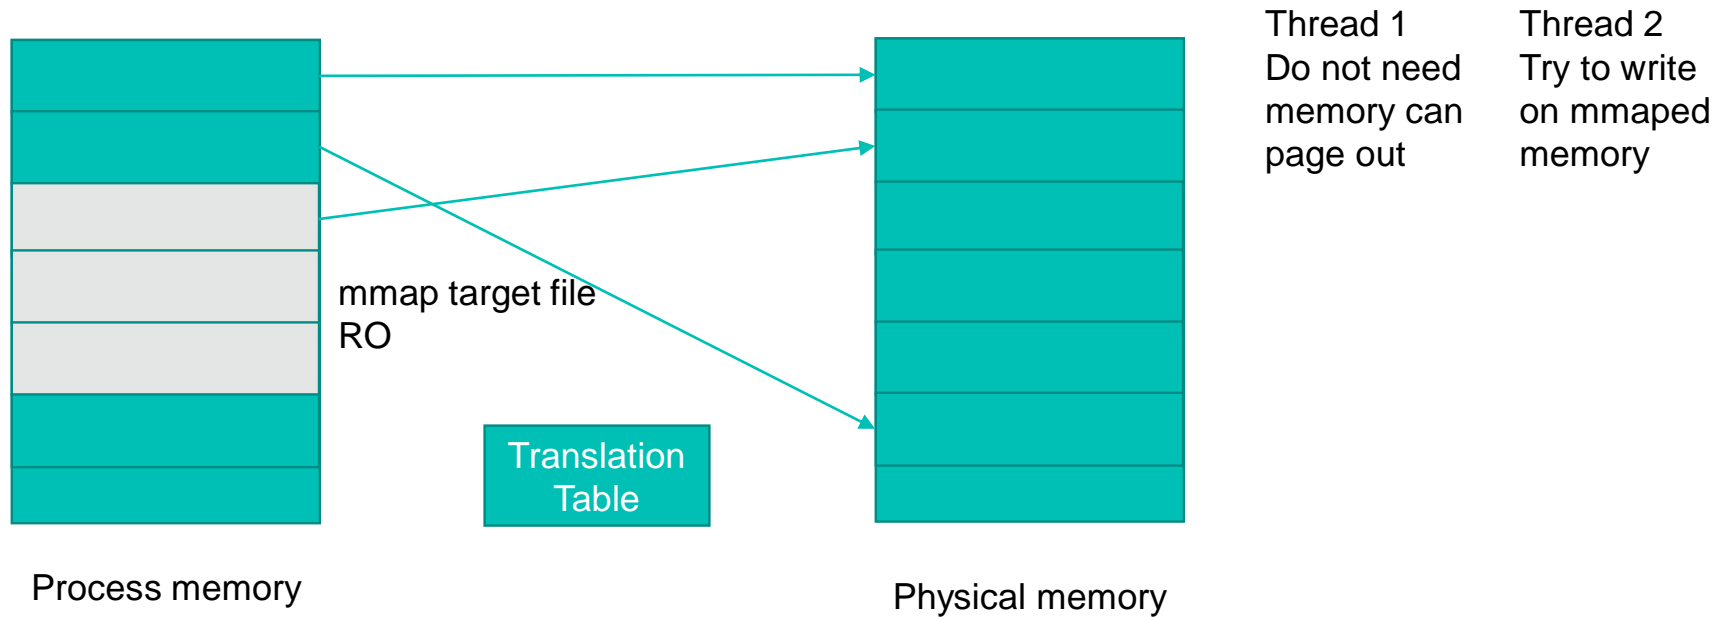
## ○ Optimization

- Two process have same value in memory
- Use the same physical address
- Fine if only READ!
- If want to read
  - COPY
  - CHANGE TT
  - WRITE
  - COW

## Example 3: dirty cow

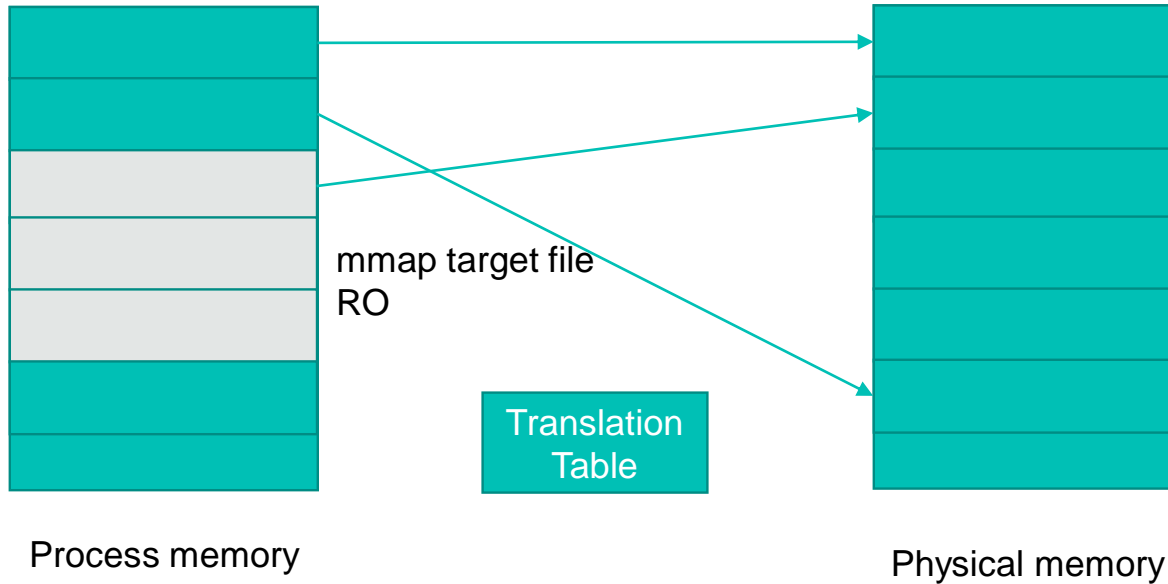


# Example 3: dirty cow





# Example 3: dirty cow




Thread 1  
Do not need  
memory can  
page out

Thread 2  
Try to write  
on mmaped  
memory

Normally fine!

# Example 3: dirty cow



Thread 1  
Do not need  
memory can  
page out

Thread 2  
Try to write  
on mmaped  
memory

Sometimes race condition  
kernel see change to mmaped file  
Copy changes to underlying file  
As it is paging out...  
... but user did not have acces!  
Can do this on...  
passwd for example

## Example 3: dirty cow

- <https://github.com/dirtycow/dirtycow.github.io/blob/master/dirtycow.c>

## Example 3: dirty cow

Homework/exam question:  
Explain the dirty cow  
vulnerability.

- <https://github.com/dirtycow/dirtycow.github.io/blob/master/dirtycow.c>

# Plan

- Race condition
- Example of race condition vulnerabilities
  - access system call
  - How (not to) implement an OS reference monitor
  - dirty cow

# Plan

Homework/exam question:  
Give an example of race  
condition vulnerability.

- Race condition
- Example of race condition vulnerabilities
  - access system call
  - How (not to) implement an OS reference monitor
  - dirty cow

# Thank you, questions?

Office MVB 3.26

[bristol.ac.uk](http://bristol.ac.uk)

